

# Haskell - Concurrency and parallelism

## Functional Programming

Jens Egholm Pedersen and Anders Kalhauge



Spring 2018

Parallelism and concurrency

Lightweight threads

STM

Data parallelism and Futhark

Definition: **code is split up and executed out-of-order**

Definition: **code is split up and executed out-of-order**

Advantage: Possible to exploit multi-core architecture

Definition: **code is split up and executed out-of-order**

Advantage: Possible to exploit multi-core architecture

Disadvantage:

1. Race-conditions
2. Deadlocking
3. Starvation

Definition: code is split up and executed *at the same time*

Definition: **code is split up and executed *at the same time***

Code can thus be *concurrent* but not *parallel*. What is an example?

- Concurrency is when many things are happening in random order
- Parallelism is when many things happen at the same time



- Concurrency is when many things are happening in random order
- Parallelism is when many things happen at the same time

- Not always *deterministic*
- Complicated to keep track of data flow
- Sometimes depending on concurrent processes (idle locks)

Correct parallelism is

- Deterministic (same outcome)
- High-level declarative (does not deal with synchronisation or communication)

Haskell provides tons of tools for this

- Par Monad
- Eval Monad
- MVars
- IO Manager
- Asynchronous exceptions
- Software Transactional Memory (STM)
- Lightweight threads

Haskell provides tons of tools for this

- Par Monad
- Eval Monad
- MVars
- IO Manager
- Asynchronous exceptions
- **Software Transactional Memory (STM)**
- **Lightweight threads**
- Lastly: Data parallelism and Futhark

A lot of the material in these slides have been borrowed from a course on parallel functional programming on Copenhagen University. Especially from slides on Haskell by Ken Friis Larsen.

Some of the material is also borrowed from the book on Parallel and Concurrent Programming in Haskell:  
<http://shop.oreilly.com/product/0636920026365.do>

In Java:

```
import java.util.Thread;  
Thread t = new Thread(() -> ...);
```

In Java:

```
import java.util.Thread;  
Thread t = new Thread(() -> ...);
```

In Haskell:

```
import Control.Concurrent  
-- forkIO :: IO () -> IO ThreadId  
  
main = do  
    threadId <- forkIO $ putStrLn "FP_ocks!"
```



Same problem as in Java: How would you get data in/out of a thread?

Same problem as in Java: How would you get data in/out of a thread?

Solution: external variables!

An MVar is exactly that: a *atomic* variable.

An MVar is exactly that: a *atomic* variable.

```
data MVar a -- abstract

newEmptyMVar :: IO (MVar a)
readMVar    :: MVar a -> IO a
takeMVar    :: MVar a -> IO a
putMVar     :: MVar a -> a -> IO ()
```

```
getURL :: String -> IO String
```

```
getURL :: String -> IO String
```

```
main = do
  m1 <- newEmptyMVar
  m2 <- newEmptyMVar
  forkIO $ do
    r <- getURL "http://www.wikipedia.org/wiki/Shovel"
    putMVar m1 r
  forkIO $ do
    r <- getURL "http://www.wikipedia.org/wiki/Spade"
    putMVar m2 r
  r1 <- takeMVar m1
  r2 <- takeMVar m2
  return (r1,r2)
```

```
data Async a = Async (MVar a)

async :: IO a -> IO (Async a)
async action = do
    var <- newEmptyMVar
    forkIO $ action >>= putMVar var
    return $ Async var

wait :: Async a -> IO a
wait (Async var) = readMVar var
```

```
main = do
  a1 <- async $
    getURL "http://www.wikipedia.org/wiki/Shovel"
  a2 <- async $
    getURL "http://www.wikipedia.org/wiki/Spade"
  r1 <- wait a1
  r2 <- wait a2
  return (r1,r2)
```



```
main = do
  a1 <- async $
    getURL "http://www.wikipedia.org/wiki/Shovel"
  a2 <- async $
    getURL "http://www.wikipedia.org/wiki/Spade"
  r1 <- wait a1
  r2 <- wait a2
  return (r1,r2)
```

```
getURLs :: [String] -> IO [ByteString]
getURLs sites =
  mapM (async.getURL) sites >>= mapM wait
```

Remember the concurrency problems:

1. Race-conditions
2. Deadlocking
3. Starvation

What is the root of these problems?

In parallel computing we generally face two problems:

- Granularity
- Data dependency

In parallel computing we generally face two problems:

- Granularity
- Data dependency

Solution: Think of parallelism as a series of transactions

Just like a SQL transaction STMs work by

1. storing the current state of affairs

Just like a SQL transaction STMs work by

1. storing the current state of affairs
2. trying to execute transaction

Just like a SQL transaction STMs work by

1. storing the current state of affairs
2. trying to execute transaction
3. roll back to previous state if error

Just like a SQL transaction STMs work by

1. storing the current state of affairs
2. trying to execute transaction
3. roll back to previous state if error
4. commit change if success



Just like a SQL transaction STMs work by

1. storing the current state of affairs
2. trying to execute transaction
3. roll back to previous state if error
4. commit change if success

Just like a SQL transaction STMs work by

1. storing the current state of affairs
2. trying to execute transaction
3. roll back to previous state if error
4. commit change if success

This is a genius use case for monads: our monad simply *is* the state

```
import Control.Concurrent.STM

data STM a -- abstract
instance Monad STM -- among other things

atomically :: STM a -> IO a
retry :: STM a
orElse :: STM a -> STM a -> STM a

data TVar a -- abstract
newTVar :: a -> STM (TVar a)
readTVar :: TVar a -> STM a
writeTVar :: TVar a -> a -> STM ()
```

```
type Amount = Int
type Account = TVar Amount
```

```
type Amount = Int
type Account = TVar Amount
```

```
transfer :: Account -> Account -> Amount -> IO ()
transfer from to amount = atomically $ do
    deposit    to    amount
    withdraw from amount
```

```
type Amount = Int
type Account = TVar Amount
```

```
transfer :: Account -> Account -> Amount -> IO ()
transfer from to amount = atomically $ do
    deposit    to    amount
    withdraw from amount
```

```
deposit :: Account -> Amount -> STM ()
deposit account amount = do
    balance <- readTVar account
    writeTVar account $ balance + amount
withdraw :: Account -> Amount -> STM ()
withdraw account amount = deposit account (- amount)
```

```
limitedWithdraw :: Account -> Amount -> STM ()
limitedWithdraw account amount = do
  balance <- readTVar account
  if amount > 0 && amount > balance
    then retry
    else writeTVar account $ balance - amount
```

```
limitedWithdraw :: Account -> Amount -> STM ()
limitedWithdraw account amount = do
    balance <- readTVar account
    if amount > 0 && amount > balance
    then retry
    else writeTVar account $ balance - amount
```

```
backupWithdraw :: Account -> Account -> Amount -> STM ()
backupWithdraw acc1 acc2 amt =
    (limitedWithdraw acc1 amt)
    'orElse '
    (limitedWithdraw acc2 amt)
```



Threads and STM still have performance problems

Threads and STM still have performance problems

What about GPU acceleration you say?

Threads and STM still have performance problems

What about GPU acceleration you say?

Functional programming is historically bad at this

Threads and STM still have performance problems

What about GPU acceleration you say?

Functional programming is historically bad at this

Solution: Data parallelism

Until now we are parallelising program logic, not data


Until now we are parallelising program logic, not data

Task parallelism: Break problem into small parts, delegate to threads

Until now we are parallelising program logic, not data

Task parallelism: Break problem into small parts, delegate to threads

Data parallelism: Break data into small parts, delegate to threads



Work: The amount of instructions to do, assuming a single PU

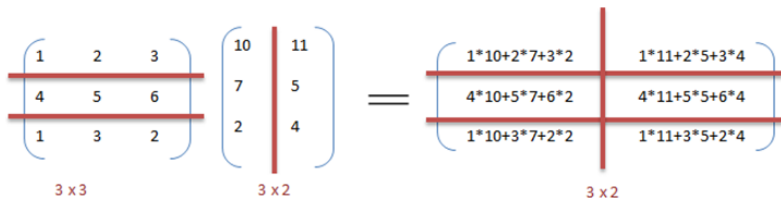
Span: The longest series of instructions, assuming infinite PU



$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} 10 & 11 \\ 7 & 5 \\ 2 & 4 \end{pmatrix} = \begin{pmatrix} 1*10+2*7+3*2 & 1*11+2*5+3*4 \\ 4*10+5*7+6*2 & 4*11+5*5+6*4 \\ 1*10+3*7+2*2 & 1*11+3*5+2*4 \end{pmatrix}$$

3 x 3                      3 x 2                      3 x 2

We can compute one cell at the time, giving  $O(n^3)$


$$\begin{array}{|c|c|c|} \hline 1 & 2 & 3 \\ \hline 4 & 5 & 6 \\ \hline 1 & 3 & 2 \\ \hline \end{array} \begin{array}{|c|c|} \hline 10 & 11 \\ \hline 7 & 5 \\ \hline 2 & 4 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 1*10+2*7+3*2 & 1*11+2*5+3*4 \\ \hline 4*10+5*7+6*2 & 4*11+5*5+6*4 \\ \hline 1*10+3*7+2*2 & 1*11+3*5+2*4 \\ \hline \end{array}$$

3 x 3                      3 x 2                      3 x 2

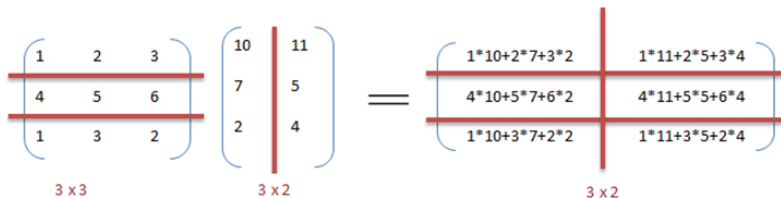
We can compute one cell at the time, giving  $O(n^3)$  (work)

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} 10 & 11 \\ 7 & 5 \\ 2 & 4 \end{pmatrix} = \begin{pmatrix} 1*10+2*7+3*2 & 1*11+2*5+3*4 \\ 4*10+5*7+6*2 & 4*11+5*5+6*4 \\ 1*10+3*7+2*2 & 1*11+3*5+2*4 \end{pmatrix}$$

3 x 3                      3 x 2                      3 x 2

We can compute one cell at the time, giving  $O(n^3)$  (work)

Or we can compute all cells in parallel, giving  $O(1)$


$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} 10 & 11 \\ 7 & 5 \\ 2 & 4 \end{pmatrix} = \begin{pmatrix} 1*10+2*7+3*2 & 1*11+2*5+3*4 \\ 4*10+5*7+6*2 & 4*11+5*5+6*4 \\ 1*10+3*7+2*2 & 1*11+3*5+2*4 \end{pmatrix}$$

$3 \times 3 \qquad 3 \times 2 \qquad 3 \times 2$

We can compute one cell at the time, giving  $O(n^3)$  (work)

Or we can compute all cells in parallel, giving  $O(1)$  (span)

This is similar to the Single Instruction Multiple Data SIMD idea

This is similar to the Single Instruction Multiple Data SIMD idea  
Incredibly useful in GPU (CUDA) programming

This is similar to the Single Instruction Multiple Data SIMD idea

Incredibly useful in GPU (CUDA) programming

Used in the OpenCL standard for GPUs

Futhark is a data parallel programming language developed here in Copenhagen

<https://futhark-lang.org/>



Futhark is a data parallel programming language developed here in Copenhagen

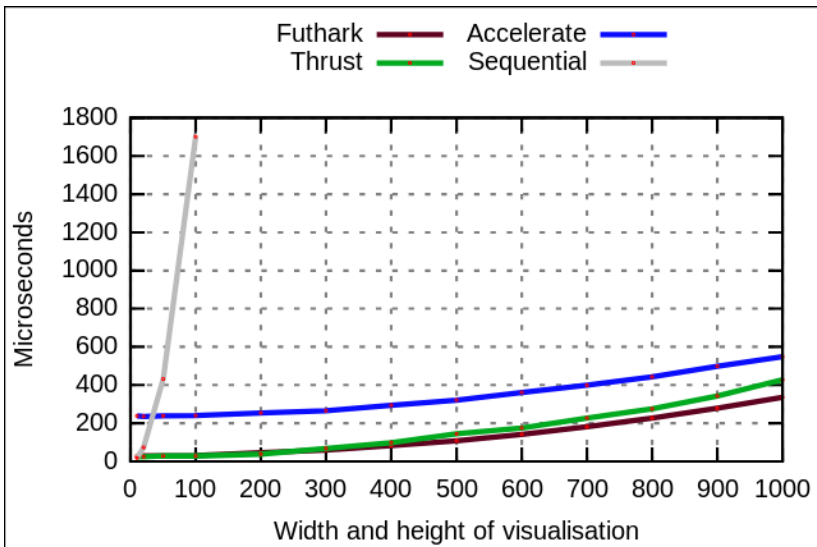
<https://futhark-lang.org/>

```
let fact (n: i32): i32 = reduce ) 1 (1...n)
```

Futhark is a data parallel programming language developed here in Copenhagen

<https://futhark-lang.org/>

```
let fact (n: i32): i32 = reduce ) 1 (1...n)
```



- ❑ Concurrency vs. parallelism
- ❑ Lightweight threading with `forkIO`
- ❑ Software Transactional Memory (STM)
- ❑ Task parallelism versus data parallelism
- ❑ Futhark and GPU accelerated functional programming